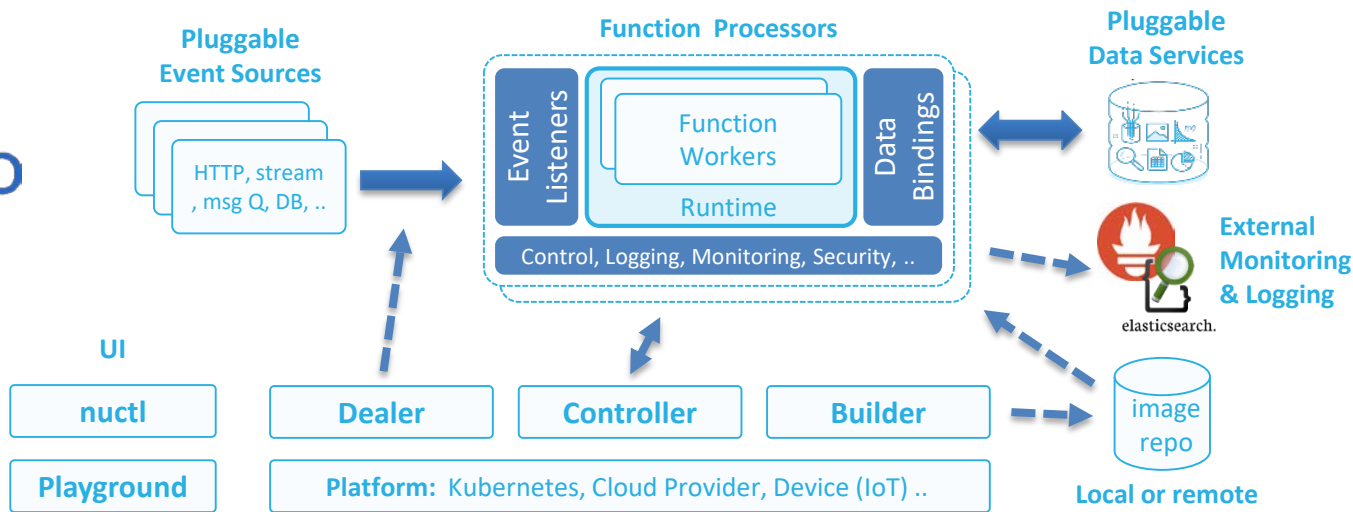


# nuclo - Comprehensive, Open, Portable, & Super Fast “serverless”

1



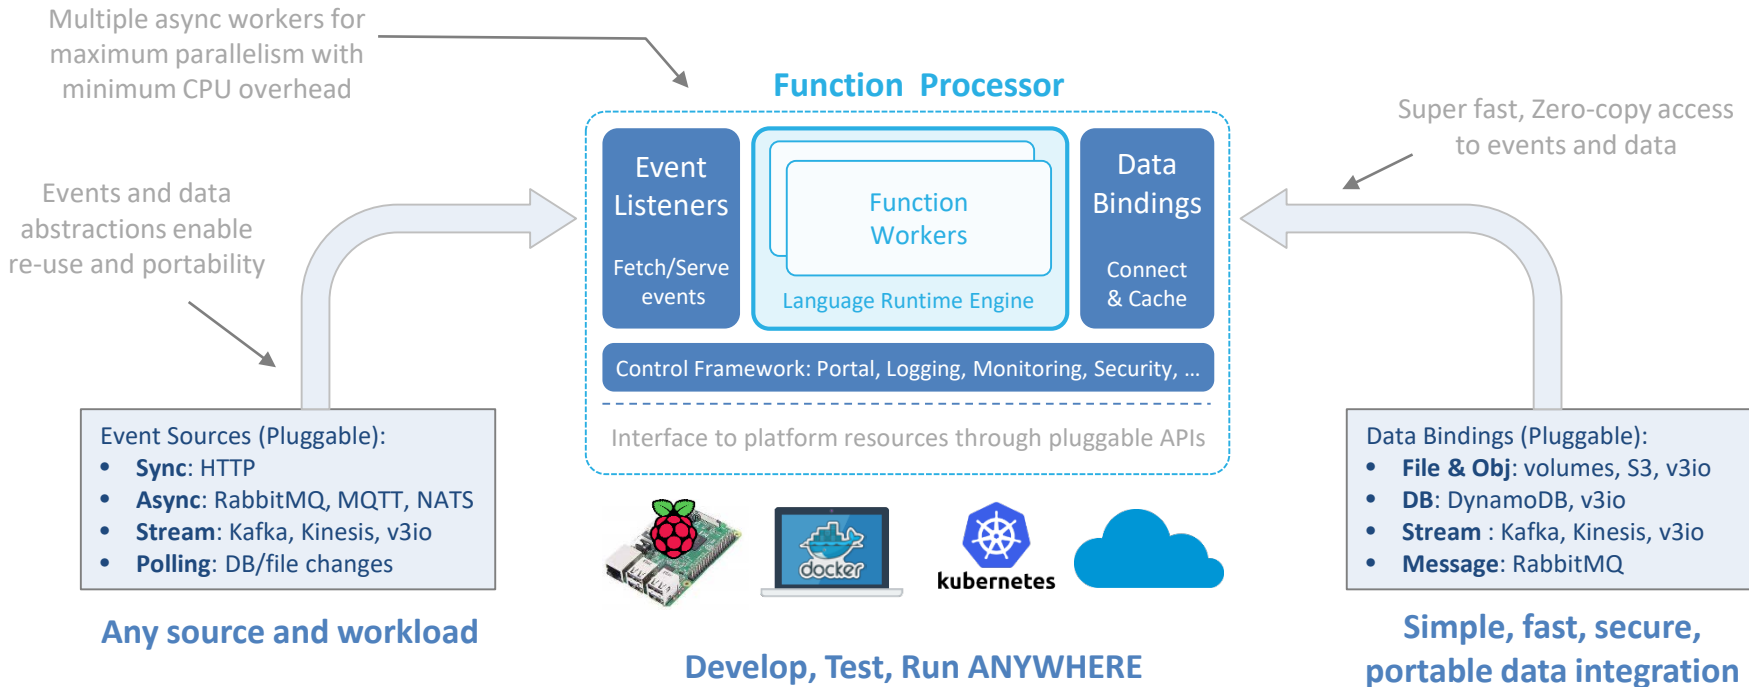
nuclo



- Real-time processing, low CPU overhead and maximum parallelism
- Simple debugging, regression, and multi-versioned CI/CD pipeline
- Pluggable data/event sources with common APIs
- Portable across low-power devices, laptops, on-prem and public cloud


# nuclo processor – Fast, Modular & Extensible

**400K events/sec per process** (100x faster than leading implementations)




# Nuclio invocation modes

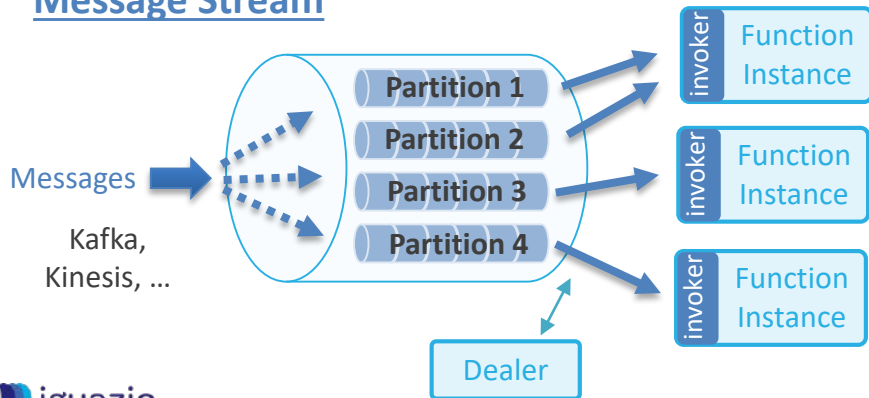
## Synchronous Req/Rep



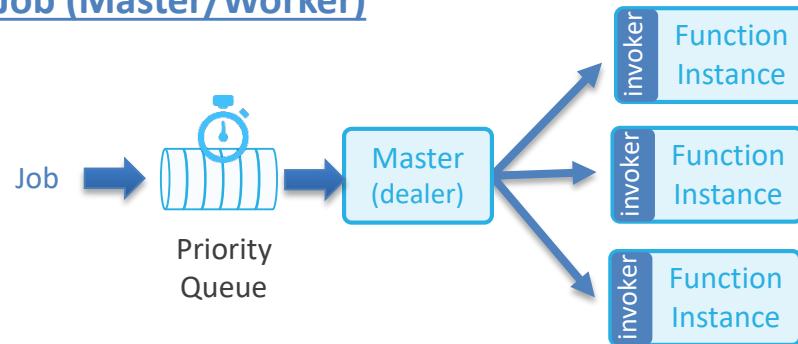
## Async Message Queue



## Message Stream




## Job (Master/Worker)

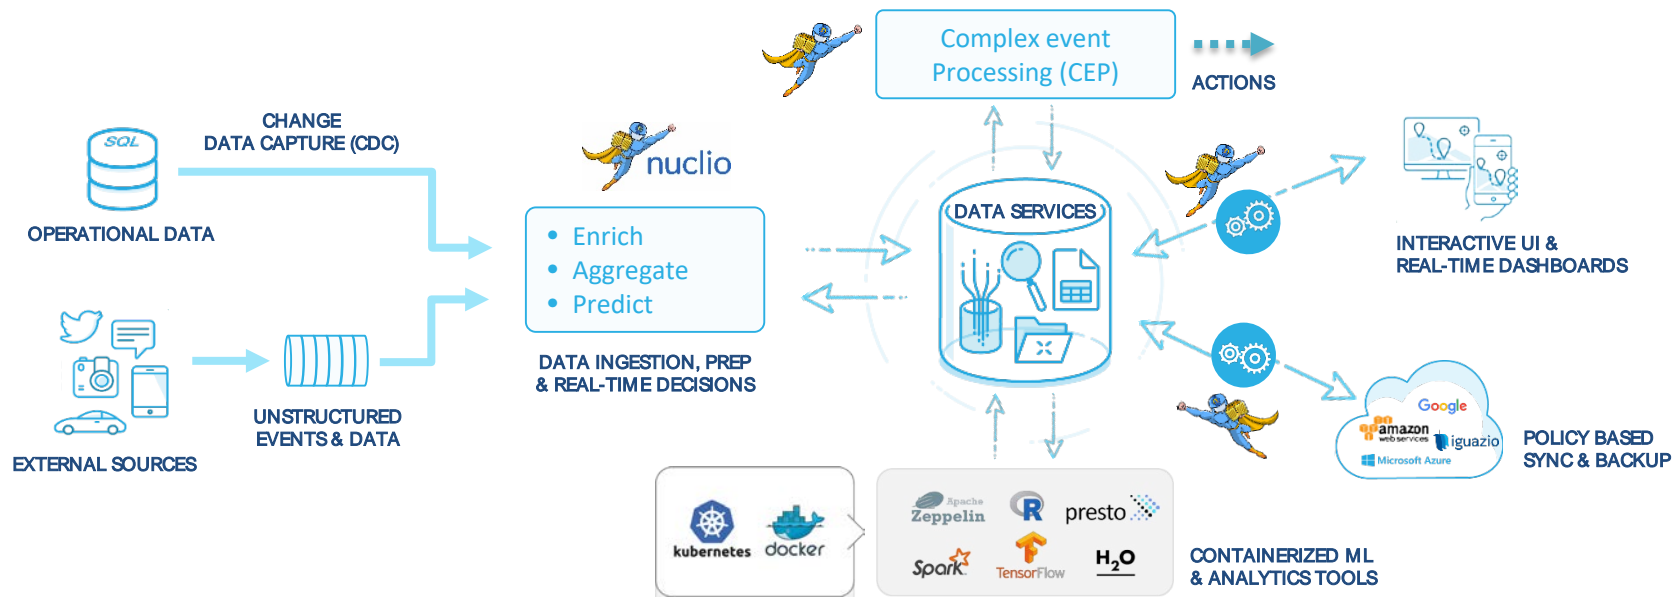


# Nuclo Dealer

- Enable real-time stream processing, batch and interactive jobs on auto-scaling Serverless functions
- By dynamically allocating tasks to workers and handling task lifecycle, checkpoint and completion




# nuclo features & Performance make Serverless broadly applicable



Higher-Productivity | Faster insights | No infrastructure hassle | Lower TCO

# Real example: Event Driven Analytics for Connected Cars

\* See code in the  
UI/Playground slide



Complex **Events** + **Data** processed in real-time without the infrastructure hassle

# nuclio

## Function Spec

**Support Kubernetes CRD:**  
Functions can be created & deleted using kubectl

```

apiVersion: "nuclio.io/v1"
kind: Function
metadata:
  name: example
  namespace: myproject
  labels:
    author: joe
spec:
  image: example:latest
  replicas: 0
  maxReplicas: 10
  env:
    - name: SOME_ENV
      value: abc
    - name: SECRET_PASSWORD_ENV_VAR
      valueFrom:
        secretKeyRef:
          name: my-secret
          key: password
  resources:
    requests:
      memory: "64Mi"
      cpu: "250m"
    limits:
      memory: "128Mi"
      cpu: "500m"
  dataBindings:
    db0:
      class: v3io
      secret: mysecret
      url: http://199.19.70.139:8081/1024
  
```

namespaced

tags/labels used for search and event sources (Label Selectors)

Various src code options\*: inline code, path (local/http/git), or local/remote pre-built image

Control Min/Max Replicas for controlled auto-scale

Pass text or secret environment variables (k8s convention)

Flex resource allocation, GPUs are coming

**Pluggable Data Sources**

\*Advanced build instructions & dependencies are in the build.yaml file

# Nuclio common event model

```

type Event interface {
    // Unique ID of the event
    GetID() string
    // Event Source class, kind, ver, schema
    GetEventSource() SourceInfoProvider
    // Event Source address (e.g. origin host IP:port, origin stream, ..)
    GetSourceAddress() string
    // Source identity (e.g. authenticated by a gateway)
    GetSourceIdentity() string
    // Content type e.g. application/json
    GetContentType() string
    // byte array of content, encoding defined by content type
    GetBody() []byte
    // Get header(s) (e.g. HTTP, AMQP, or anything injected by the source)
    // also have convenience methods: GetHeaderByteSlice, GetHeaderString
    GetHeader(key string) interface{}
    GetHeaders() map[string]interface{}
    // Get field(s), decode fields in the body (e.g. json, DB record, ..)
    // Allow functions to ignore the specific event encoding, e.g. emulate DB record via HTTP
    // also have convenience methods: GetFieldByteSlice, GetFieldString, GetFieldInt
    GetField(key string) interface{}
    GetFields() map[string]interface{}
    // Original event timestamp or gateway timestamp (if origin timestamp not specified)
    GetTimestamp() time.Time
    // Logical path requested by the event (e.g. HTTP request path, stream name, etc.)
    GetPath() string
    // URL object for HTTP requests, for convenience
    GetURL() URL
    // HTTP or transport method
    GetMethod() string
    // Translate event to a json byte array
    AsJson() []byte
}

```

Simplify and generalize  
client implementation

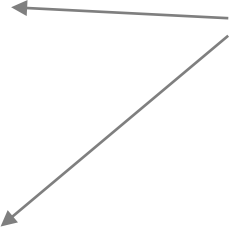
Enable zero copy and zero  
ser/des when possible



# Context.logger Interface

```
type Logger interface {  
  
    // emit a log entry of a given verbosity. the first argument may be an object, a string  
    // or a format string. in case of the latter, the following varargs are passed  
    // to a formatter (e.g. fmt.Sprintf)  
    Error(format interface{}, vars ...interface{})  
    Warn(format interface{}, vars ...interface{})  
    Info(format interface{}, vars ...interface{})  
    Debug(format interface{}, vars ...interface{})  
  
    // emit a structured log entry. example:  
    //  
    // l.InfoWith("The message",  
    //           "first-key", "first-value",  
    //           "second-key", 2)  
    //  
    ErrorWith(format interface{}, vars ...interface{})  
    WarnWith(format interface{}, vars ...interface{})  
    InfoWith(format interface{}, vars ...interface{})  
    DebugWith(format interface{}, vars ...interface{})  
  
    // flushes buffered logs, if applicable  
    Flush()  
  
    // returns a child logger, if underlying logger supports hierarchal logging  
    GetChild(name string) interface{}  
}
```

Support both structured & unstructured logging



Support nested/hierarchical logs




One log interface, multiple implementations (screen, file, stream, http, ..), extensible

# Default Context.DataBinding API (sync & async ver), can be overwritten

Service	Major APIs	Main Request Params
<b>Object</b> e.g. S3, Minio, v3io	<b>ListObjects</b> <b>GetObject</b> <b>PutObject</b> <b>DeleteObject</b>	Bucket, Prefix, MaxKeys Bucket, Key, Range Bucket, Key, Metadata, Body Bucket, Key
<b>NoSQL</b> e.g. DynamoDB, Cassandra, v3io	<b>GetItem</b> <b>GetItems</b> <b>PutItem</b> <b>UpdateItem</b> <b>DeleteItem</b>	Table, Key, Projection Table, ConditionExpression, ProjectionExpression, Limit Table, Key, ProjectionExpression, item Table, Key, UpdateExpression, ConditionExpression Table, Key, ConditionExpression
<b>Stream</b> e.g. Kinesis, Kafka, v3io	<b>GetRecords</b> <b>PutRecords</b> <b>Seek</b>	Stream, ShardId, Location, Limit Stream, Records Stream, ShardId, SeekType, SeekTime, StartingSequence, Timestamp
<b>File</b>	<b>Open</b> <b>Read</b> <b>Write</b>	Path, Mode, flags Handle, offset, size Handle, offset, size, data

# Nuclo Playground (run as isolated k8s deployment)


nuclo

Deploy

```

34
35 func CarEvent(context 'nuclo.Context, event nuclo.Event) (interface{}, error) {
36     cont := context.DataBinding["db0"]
37
38     record := carEvent{}
39     err := json.Unmarshal(event.GetBody(), &record)
40     if err != nil {
41         return nil, fmt.Errorf("Failed to Unmarshal json")
42     }
43     geoHash := GetGeoHash(record.Lon, record.Lat)
44
45     dc := dataframe.NewDataContext(context.Logger)
46
47     // Async (parallel) loading of cars and road info data
48     cars := dc.Read.FromTable(cont, "cars").Keys(record.VIN).LoadAsync()
49     roadMap := dc.Read.FromTable(cont, "roadinfo").Keys(geoHash).Select("weather", "speedlimit", "trafficCondition").LoadAsync()
50     car := cars.Next()
51     road := roadMap.Next()
52
53     if cars.Error() != nil || car==nil {
54         return nil, fmt.Errorf("Car %s not found", record.VIN)
55     }
56
57     // Sync (blocking) read of driver data, single row get by key (DriverID)
58     driver := dc.Read.FromTable(cont, "driver").Keys(car["DriverID"]).GetRow()
59
60     // Write Car, Road, and Driver data to the log at Info level
61     context.Logger.InfofWith("Got data", "car", car, "road", road, "driver", driver)
62
63     // Create an enriched data record to be streamed into various queues
64     enriched := enrichedEvent{carEvent:record, GeoHash:geoHash}
65     car.Scan("DriverID,state,lastGeoHash",&enriched.DriverID, &enriched.State, &enriched.LastGeoHash)
66     road.Scan("weather,trafficCondition",&enriched.Weather, &enriched.TrafficCondition)
67     enrichedMsg, err := json.Marshal(enriched)
68
69     // Write enriched data into stream (to be used by Spark Streaming for Machine Learning)
70     dc.Write.ToStream(cont, "enriched-stream").Records(enrichedMsg).SaveAsync()
71
72     // If car changed location update the geo map counters (-1 in old location, +1 in new location)
73     if enriched.LastGeoHash != geoHash {
  
```

Configure

Invoke

POST
Send

Content type: JSON
Log level: Info

```
{
  "VIN": "1",
  "Speed": 19,
  "State": "ok",
  "Lat": "40.7513890",
  "Lon": "-73.9930560"
}
```

Log

> Body:

```
{
  "VIN": "1",
  "Lon": "-73.9930560",
  "Lat": "40.7513890",
  "Speed": 19,
  "State": "ok",
  ...
}
```

Clear log

# CLI (run command example)

```
$ nuctl run --help
Build, deploy and run a function
```

```
Usage:
  nuctl run function-name [flags]
```

## Flags:

--data string	Comma separated list of data bindings (in json)
--data-bindings string	JSON encoded data bindings for the function
--desc string	Function description
-d, --disabled	Start function disabled (don't run yet)
-e, --env string	Environment variables (name1=val1,name2=val2..)
--events string	Comma separated list of event sources (in json)
-f, --file string	Function Spec File
-h, --help	help for run
-i, --image string	Docker image name, will use function name if not specified
-l, --labels string	Additional function labels (lbl1=val1, lbl2=val2..)
--max-replica int32	Maximum number of function replicas
--min-replica int32	Minimum number of function replicas
--no-pull	Don't pull base images - use local versions
--nuclio-src-dir string	Local directory with nuclio sources (avoid cloning)
--nuclio-src-url string	nuclio sources url for git clone (default "https://github.com/nuclio/nuclio.git")
-o, --output string	Build output type - docker binary (default "docker")
-p, --path string	Function source code path
--port int32	Public HTTP port (node port)
--publish	Publish the function
-r, --registry string	URL of container registry (env: NUCTL_REGISTRY)
--run-registry string	The registry URL to pull the image from, if differs from -r (env: NUCTL_RUN_REGISTRY)
--runtime string	Runtime - golang, python, ..
-s, --scale string	Function scaling (auto number) (default "1")
--version string	Docker image version (default "latest")

## Global Flags:

-k, --kubeconfig string	Path to Kubernetes config (admin.conf) (default ~/.kube/config)
-n, --namespace string	Kubernetes namespace (default "default")
-v, --verbose	verbose output

# Perf results, single process, using basic function

```

2017/09/28 17:08:01 Starting http blaster
2017/09/28 17:08:01 Running test on [1:38573, tls node=false, block size=5, test timeout 2m40s
2017/09/28 17:08:01 Adding executor for get
2017/09/28 17:08:01 at executor start {multi get from 28 workers GET 1m0s 0 128 0 map[] 0 0 false performance 0
0 0 }
2017/09/28 17:08:01 Wait for executors to finish
2017/09/28 17:09:01 Ending multi get from 28 workers
2017/09/28 17:09:01 report for wL 0 :
2017/09/28 17:09:01 Total Requests 23417291
2017/09/28 17:09:01 Min: 0s
2017/09/28 17:09:01 Max: 211.343548ms
2017/09/28 17:09:01 Avg: 1.807384ms
2017/09/28 17:09:01 Error Count: 0
2017/09/28 17:09:01 Statuses:
2017/09/28 17:09:01 503 - 19
2017/09/28 17:09:01 200 - 23417272
2017/09/28 17:09:01 lops: 390288
2017/09/28 17:09:01 status code 200 occurred 99.999919% during the test "multi get from 28 workers"
2017/09/28 17:09:01 report for wL 0 :
2017/09/28 17:09:01 Total Requests 23417291
2017/09/28 17:09:01 Min: 0s
2017/09/28 17:09:01 Max: 211.343548ms
2017/09/28 17:09:01 Avg: 1.807384ms
2017/09/28 17:09:01 Error Count: 0
2017/09/28 17:09:01 Statuses:
2017/09/28 17:09:01 200 - 23417272
2017/09/28 17:09:01 503 - 19
2017/09/28 17:09:01 lops: 390288
2017/09/28 17:09:01 Duration: 1m0.476341776s
2017/09/28 17:09:01 Overall Results:
2017/09/28 17:09:01 Overall Requests: 23417291
2017/09/28 17:09:01 Overall GET Requests: 23417291
2017/09/28 17:09:01 Overall GET Min Latency: 0s
2017/09/28 17:09:01 Overall GET Max Latency: 211.343548ms
2017/09/28 17:09:01 Overall GET Avg Latency: 1.807384ms
2017/09/28 17:09:01 Overall PUT Requests: 0
2017/09/28 17:09:01 Overall PUT Min Latency: 0s
2017/09/28 17:09:01 Overall PUT Max Latency: 0s
2017/09/28 17:09:01 Overall PUT Avg Latency: 0s
2017/09/28 17:09:01 Overall IOPS: 390288
2017/09/28 17:09:01 Overall GET IOPS: 390288
2017/09/28 17:09:01 Overall PUT IOPS: 0

```

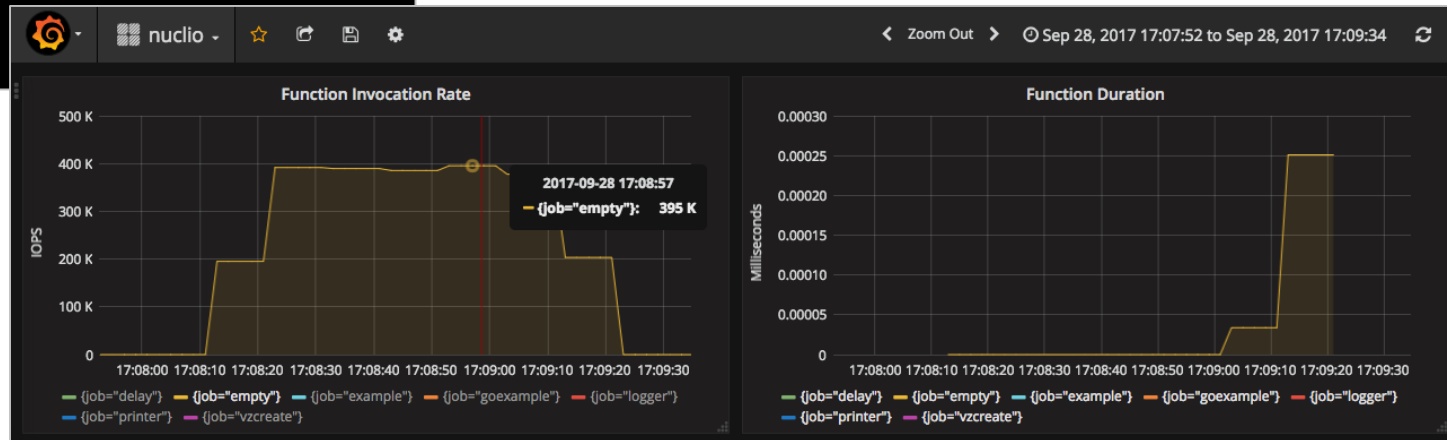
```
package empty
```

```
import (
    "github.com/nuclio/nuclio-sdk"
)
```


```
func Empty(context *nuclio.Context, event nuclio.Event) (interface{}, error) {
    return nil, nil
}
```

Tested using:

[https://github.com/v3io/http\\_blaster](https://github.com/v3io/http_blaster)



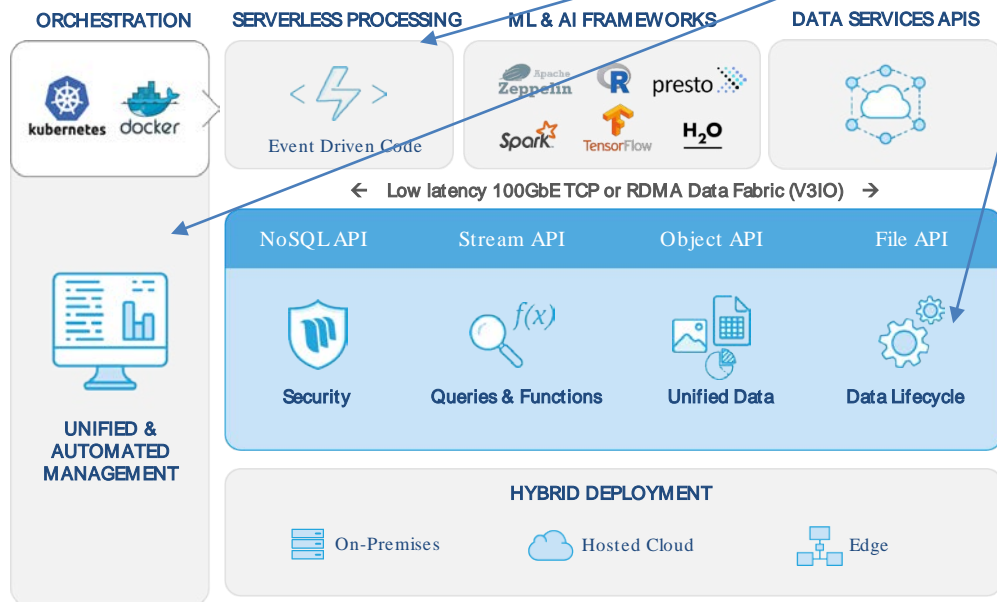
# Enabling Simplest and Continuous Dev & Ops (CI/CD)



```
$ nuctl run <name> <source> [options]
```

One Click to test, deploy, upgrade or rollback code  
Runs ANYWHERE, Self-healing and Auto-Scaling





- **Used in iguazio platform**
  - Developed for the real world
- **Now completely re-written to:**
  - Support the broader open source & CNCF eco-system
  - Incorporate learnings from G1
  - Future proof the architecture
  - Address new use cases